



The artifact-to-claim gradient, and the denominators it depends on

13 scored rows over 12 distinct packages. Source: E1-DATASET-v1.0.md (frozen 2026-09-29, corrected 2026-09-30). Every quantity below is recomputed from the frozen row table by figures/make_figures.py, which aborts if any value disagrees.

(a) Categorical layer verdicts, n = 13



(b) Adjudicated claim pool, 375 claims



Granularity conditioning

The 70.1% correct rate is computed over a pool of 375 claims that mixes two counting conventions. Batches 1-2: $213/270 = 78.9\%$. Batch 3: $28/81 = 34.6\%$ over four packages, $50/105 = 47.6\%$ over five rows. The split is a property of the splitter, so the pooled rate is descriptive and is not a defect rate for this corpus. Only panel (a) is comparable across batches.

Verdict vs. containment

Package-level verdict: 0 SUPPORTED, 8 PARTIALLY, 5 CONTRADICTED, of 13. Packages containing at least one CLAIM-CONTRADICTED claim: 13 of 13. The second is the stronger statement and the more damaging one, and it holds for the package that is SUPPORTED on every layer below its verdict. Reporting only the first understates the finding; reporting the second as the verdict misdescribes the adjudication.

Aggregate reconciliation

Two arithmetic errors in the frozen dataset's derived-metrics block, found by row-level reconciliation against the thirteen per-package rows and corrected with the rows left untouched. U4 verdict distribution: block recorded 0 / 7 / 6, corrected to 0 / 8 / 5. internal consistency: block recorded 2 of 13, corrected to 1 of 13. The correction makes the gradient steeper, not flatter: 1 of 13 rows are internally consistent. Neither correction changes the central result: 0 of 13 packages are SUPPORTED and 13 of 13 contain a contradicted claim before and after. Section 4.4 gives both rounds.